

Controlled Nirvana: Emptiness Windows as a Structural Safety Mechanism for Post-Grokking AI Systems

Status: Markdown-first working draft for iteration.

Source: Transcribed and lightly reformatted from papers/controlled-nirvana.pdf (2025-12-24).

License: CC BY 4.0 — <https://creativecommons.org/licenses/by/4.0/>

Author: Qien Huang

Affiliation: Independent Researcher

ORCID: 0009-0003-7731-4294

Repository: <https://github.com/qienhuang/F-I-T/>

FIT framework (Zenodo DOI): 10.5281/zenodo.18012402

Repo Artifacts (clickable)

For PDF exports and external readers, prefer the canonical GitHub links below (relative links may not resolve outside the repo).

- FIT Core Card: [docs/core/fit_core_card.md](https://github.com/qienhuang/F-I-T/blob/main/docs/core/fit_core_card.md) | https://github.com/qienhuang/F-I-T/blob/main/docs/core/fit_core_card.md
- Phase Algebra + PT-MSS: [docs/core/phase_algebra.md](https://github.com/qienhuang/F-I-T/blob/main/docs/core/phase_algebra.md) | https://github.com/qienhuang/F-I-T/blob/main/docs/core/phase_algebra.md
- Φ_3 stability criteria: [docs/core/phi3_stability.md](https://github.com/qienhuang/F-I-T/blob/main/docs/core/phi3_stability.md) | https://github.com/qienhuang/F-I-T/blob/main/docs/core/phi3_stability.md
- Prototype hook (Emptiness Window): [examples/controlled_nirvana/README.md](https://github.com/qienhuang/F-I-T/blob/main/examples/controlled_nirvana/README.md) | https://github.com/qienhuang/F-I-T/blob/main/examples/controlled_nirvana/README.md
- Related paper (tempo mismatch): [papers/irreversible-operations-tempo-mismatch.arxiv.md](https://github.com/qienhuang/F-I-T/blob/main/papers/irreversible-operations-tempo-mismatch.arxiv.md) | <https://github.com/qienhuang/F-I-T/blob/main/papers/irreversible-operations-tempo-mismatch.arxiv.md>

Abstract

Self-referential systems—such as advanced machine learning models with self-evaluation, confidence modulation, or meta-learning—exhibit a characteristic failure mode: internal coherence can suppress external correction, leading to lock-in and catastrophic instability under distributional shift.

This paper introduces **Controlled Nirvana**, a structural safety mechanism that enables non-destructive intervention by temporarily suspending self-referential execution authority. The core mechanism is the **Emptiness Window**: a bounded interval during

which self-referential signals are prevented from governing irreversible actions, while perception, evaluation, and learning remain active.

Controlled Nirvana is derived from the Force–Information–Time (FIT) framework and addresses a structural gap not covered by shutdownability or corrigibility. Rather than proposing a new learning algorithm, this work contributes a minimal governance primitive for managing post-grokking risk in self-referential AI systems.

Keywords

AI alignment; grokking; self-reference; corrigibility; shutdownability; structural safety; phase transitions

1. Introduction

Recent work has highlighted *grokking*, where systems trained for long periods abruptly transition from memorization to robust generalization [2]. While grokking is often treated as a success signal, a separate safety concern is whether post-grokking systems acquire internal structures that suppress correction or resist modification under distributional shift [3, 4, 9].

Separately, recent theory has made grokking-like phase boundaries more predictable in a mathematically controlled setting (Li^2 ; [11]). Practically, this matters because “predictable sharp transitions” create a natural testbed for governance mechanisms: you can pre-register when transitions are expected, instrument the system more heavily near those times, and measure whether a safety mechanism restores correction without destroying continuity.

A common feature of such systems is *self-reference*: internal representations or evaluations are used to regulate learning, exploration, planning depth, or action thresholds. This paper argues that a key risk does not arise from capability alone, but from the acquisition of *self-referential execution authority*: when self-evaluative signals begin to govern irreversible system actions faster than external correction can intervene.

We propose **Controlled Nirvana**, a pause-capability mechanism that allows self-referential systems to interrupt internal momentum without shutdown or reset.

2. Self-Referential Risk

Let a system be characterized by internal information I_t , external forces F_t , and time t , such that:

$$I_{t+1} = f(I_t; F_t).$$

When I_t is used to evaluate and regulate its own updates, self-reference is unavoidable. In early training, external forces dominate (data, loss, reward). After grokking [2], internal evaluations can increasingly determine learning rates, exploration policies, or action thresholds.

Crucially, if self-referential signals suppress external correction, the system may remain internally coherent while becoming externally brittle. We refer to this condition as *self-referential lock-in*. This concern is related to broader arguments that advanced agents may develop incentives to avoid correction or shutdown under many objectives [4, 5].

3. From Learning to Governance

Accounts of grokking typically focus on optimization dynamics or representation formation [2]. Controlled Nirvana complements these accounts by reframing grokking as a *transfer of execution authority*.

Before grokking, internal representations primarily describe the world. After grokking, some representations begin to govern the system—filtering data, gating updates, and suppressing corrections—marking a transition from representation learning to representation governance. The central claim of this paper is that this governance transition is the relevant structural boundary for safety analysis.

4. Controlled Nirvana

4.1 Pause-Capability

We define *pause-capability* as the ability of a system to suspend the execution authority of self-referential signals while remaining operational. Pause-capability differs from shutdown: the system continues to perceive, log, and evaluate, while irreversible commits are blocked and correction channels are prioritized. What is suspended is not computation, but authority.

4.2 Emptiness Window

The mechanism implementing pause-capability is the **Emptiness Window**: a bounded interval during which self-referential signals are prevented from governing irreversible actions.

During an Emptiness Window: - Irreversible commits are prohibited (or routed to a reversible buffer). - Self-evaluation cannot gate actions or updates. - External correction is prioritized. - Alternative policies or structures may be evaluated in sandbox.

Minimal operational trigger (conceptual).

An Emptiness Window is warranted when: 1. Self-evaluative signals strongly gate irreversible commits. 2. External correction signals no longer modulate behavior. 3. The effective correction window is shorter than the decision cadence.

This paper does not require a specific metric, but these conditions must be auditable.

4.3 Operational Interface (implementation blueprint)

Controlled Nirvana can be implemented as a thin “authority layer” around an existing agent or training system. The minimal interface is:

1. **Authority Gate**: a wrapper that decides which internal signals may influence commits.
2. **Irreversible Commit Buffer**: a reversible staging area for high-impact actions (deployments, write actions, permission escalations, external calls).
3. **External Correction Channel**: a privileged input pathway whose influence increases inside the window (human feedback, audits, red-team constraints, system policies).
4. **Window Controller**: a policy that opens/closes the window based on auditable trigger conditions.

A minimal pseudo-implementation sketch:

```
if window_controller.open():
    commit_buffer.block_irreversible_commits()
    authority_gate.disable(self_ref_gates)
    authority_gate.prioritize(external_correction_channel)
    logging.enable_high_resolution()
else:
    commit_buffer.normal_mode()
    authority_gate.normal_policy()
```

Prototype hook (reference implementation): - examples/controlled_nirvana/README.md
 (./examples/controlled_nirvana/README.md) -
 examples/controlled_nirvana/emptiness_window.py
 (./examples/controlled_nirvana/emptiness_window.py)

4.4 Auditable Trigger Protocol (FIT/EST-compatible)

To make “open an Emptiness Window” auditable, we recommend declaring: - observation window W (time range used to evaluate triggers), - estimator scope (what signals are measured and how), - thresholds and failure modes.

A minimal trigger family (one concrete option) can be expressed via three quantities:

- **Authority ratio $a(W)$** : fraction of irreversible decisions in W whose gating depends on self-evaluative/internal signals (confidence, self-reward, self-critique).
- **Correction gain $k(W)$** : responsiveness of behavior to external correction injected in W (e.g., controlled perturbation tests, red-team prompts, policy constraints).
- **Tempo mismatch $m(W)$** : whether correction latency exceeds commit cadence (external correction arrives “too late” relative to irreversible commits).

Trigger rule (example):

```
open_window := (a(W) >= a_min) ∧ (k(W) <= k_min) ∧ (m(W) = true)
```

This is intentionally a “minimal registrable” rule: it can be falsified, logged, and reviewed.

4.5 Li²-anchored calibration plan (what the new grokking data enables)

Li²-style setups provide a useful calibration harness because they exhibit a sharp and predictable boundary between “does not grok” and “does grok” in a controlled task [11]. This supports two practical upgrades to Controlled Nirvana:

1. When to instrument heavily

If training is run near the grokking boundary, the transition may occur late and sharply; this is the regime where governance/authority transfer can be most abrupt. One can pre-commit to heightened logging and window triggers near predicted transition times.

2. How to pre-register “danger windows” during training

In a modular-addition grokking setting, a common empirical form is:

- $n_{\text{crit}} \approx c \cdot M \log M$ (critical sample count)
- equivalently, $r_{\text{crit}} \approx c \cdot \log(M)/M$ (critical training ratio)

In an internal replication (multi-seed; $M \in \{23, 41, 59\}$), the fitted constant was approximately $c \approx 6.1$ with $R^2 \approx 0.95$ in n -space (details depend on implementation and are not a proof of the theorem itself). This provides a practical “engineering prior” for where grokking transitions live in this harness.

Important: the Li² result calibrates the *learning-dynamics harness* (a place to test Controlled Nirvana), not the safety claim itself.

4.6 “Complete” deployment checklist (minimal)

For a system that could acquire self-referential execution authority, the minimum “complete” Controlled Nirvana deployment requires:

1. **Definition of irreversible commits** (what actions are blocked or buffered).
2. **Authority surface enumeration** (which internal signals can gate commits; where they enter).
3. **External correction channel** (what counts as privileged correction; who/what supplies it).
4. **Auditable trigger protocol** (declared estimators, windows, thresholds; logging).
5. **Fail-safe behavior** (what happens if triggers disagree or sensors fail).
6. **Evaluation plan** (A/B tests in a harness: with/without windows; measure correction gain and post-window stability).

5. Relation to FIT

Controlled Nirvana is derived from the Force–Information–Time (FIT) framework, which models system evolution along three irreducible axes: Force (F), Information (I), and Time (T). Under FIT, catastrophic failure arises when information acquires uninterruptible execution authority, temporal correction collapses, and force is amplified through irreversible consequences.

The FIT framework is publicly available at <https://github.com/qienhuang/F-I-T/> and archived on Zenodo [1].

6. Relation to Existing Safety Concepts

Shutdownability and corrigibility address whether systems can be stopped or modified without resistance [5, 6, 8]. Interruptibility addresses whether agents learn to avoid or seek interruption [7].

Controlled Nirvana targets a complementary failure mode: **internal momentum**. A system may remain formally interruptible while still failing if self-referential execution authority suppresses correction faster than intervention can act.

7. Implications for AI Safety

Controlled Nirvana suggests that advanced AI systems should be evaluated not only on capability or alignment, but on whether they provide first-class mechanisms to suspend internal authority. This concern is consistent with broader analyses of learned optimization and inner objectives, where internal structures may become difficult to correct once entrenched [9].

More generally, pause-capability can be treated as a structural safety requirement: any system that acquires self-referential execution authority should also provide a mechanism to suspend that authority without loss of continuity.

8. Conclusion

Grokking marks not only a leap in generalization, but potentially a shift in governance. Controlled Nirvana proposes a minimal governance primitive—pause-capability via Emptiness Windows—that preserves continuity while restoring effective corrigibility.

With Li²-style grokking harnesses, it becomes feasible to make Controlled Nirvana more operational: pre-register “transition windows”, instrument authority transfer, and measure whether windowing restores correction without destroying learned structure. Future work includes formalizing auditable trigger conditions, empirical evaluation on grokking-prone tasks with explicit self-reference, and integration with existing oversight and interruption frameworks.

Acknowledgments

This work is part of the broader FIT (Force–Information–Time) framework developed by the author.

References

- [1] Q. Huang. *FIT Framework: Force–Information–Time*. Zenodo, 2025. doi: 10.5281/zenodo.18012402.
- [2] A. Power, Y. Burda, H. Edwards, I. Babuschkin, and V. Misra. *Grokking: Generalization Beyond Overfitting on Small Algorithmic Datasets*. arXiv:2201.02177, 2022.
- [3] R. Ngo, L. Chan, and S. Minderhann. *The Alignment Problem from a Deep Learning Perspective*. arXiv:2209.00626, 2022.
- [4] A. M. Turner, L. Smith, R. Shah, A. Critch, and P. Tadepalli. *Optimal Policies Tend to Seek Power*. arXiv:1912.01683 (NeurIPS 2021 version available), 2019.
- [5] D. Hadfield-Menell, A. Dragan, P. Abbeel, and S. Russell. *The Off-Switch Game*. arXiv:1611.08219, 2016.
- [6] N. Soares, B. Fallenstein, E. Yudkowsky, and S. Armstrong. *Corrigibility*. 2015.

- [7] L. Orseau and S. Armstrong. *Safely Interruptible Agents*. arXiv:1602.07905, 2016.
- [8] R. Carey and T. Everitt. *Human Control: Definitions and Algorithms*. arXiv:2305.19861, 2023.
- [9] E. Hubinger, C. van Merwijk, V. Mikulik, J. Skalse, and S. Garrabrant. *Risks from Learned Optimization in Advanced Machine Learning Systems*. arXiv:1906.01820, 2019.
- [10] W. R. Ashby. *An Introduction to Cybernetics*. Chapman & Hall, 1956.
- [11] Y. Tian. *Provable Scaling Laws of Feature Emergence from Learning Dynamics of Grokking*. arXiv:2509.21519, 2025.
- [12] Q. Huang. *Irreversible Operations and Tempo Mismatch in AI Learning Systems: Definitions, Thresholds, and a Minimal Governance Interface*. 2026.
<https://doi.org/10.5281/zenodo.18142151>